

Human Activity Recognition Project

Roberto Preste

2018-10-24

Overview

One thing that people regularly do is quantify *how much* of a particular activity they do, but they rarely quantify *how well* they do it. The goal of this project will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants who were asked to perform barbell lifts correctly and incorrectly in 5 different ways, and predict the manner in which they did the exercise. More information can be found on this web page.

Loading data

```
library(tidyverse)
library(caret)
```

Let's load the data into R.

```
training <- read.csv("pml-training.csv")
dim(training)
```

```
## [1] 19622 160
```

```
testing <- read.csv("pml-testing.csv")
dim(testing)
```

```
## [1] 20 160
```


Saving the target variable from the `training` dataset and the problem ID from the `testing` dataset.

```
target <- training[, "classe"]
prob_ids <- testing[, "problem_id"]
```

Let's see if our classes are balanced.

```
plot(target, main = "Target feature distribution",
      xlab = "classe", ylab = "Count", col = "dodgerblue")
```

Target feature distribution



Data cleaning

Since the assignment requires us to use data from accelerometers on the belt, forearm, arm and dumbbell, let's first select features related to them.

```
feats <- grepl("arm|belt|dumbbell", names(training))
training <- training[, feats]
dim(training)
```

```
## [1] 19622 114
```

```
testing <- testing[, feats]
dim(testing)
```

```
## [1] 20 114
```

We have removed 46 columns, now we can drop all features having NA values, based on features in the testing dataset.

```
nonNA <- colSums(is.na(testing)) == 0
training <- training[, nonNA]
dim(training)
```

```
## [1] 19622 39
```

```
testing <- testing[, nonNA]
dim(testing)
```

```
## [1] 20 39
```

We kept 39 of the starting 160 features. This might not be the proper approach in some cases, but it should be alright for this project. As a final check, let's look for uninformative features, namely those having zero

or near-zero variance.

```
nearZeroVar(training, saveMetrics = T)
```

##	freqRatio	percentUnique	zeroVar	nzv
## roll_belt	1.101904	6.7781062	FALSE	FALSE
## pitch_belt	1.036082	9.3772296	FALSE	FALSE
## yaw_belt	1.058480	9.9734991	FALSE	FALSE
## total_accel_belt	1.063160	0.1477933	FALSE	FALSE
## gyros_belt_x	1.058651	0.7134849	FALSE	FALSE
## gyros_belt_y	1.144000	0.3516461	FALSE	FALSE
## gyros_belt_z	1.066214	0.8612782	FALSE	FALSE
## accel_belt_x	1.055412	0.8357966	FALSE	FALSE
## accel_belt_y	1.113725	0.7287738	FALSE	FALSE
## accel_belt_z	1.078767	1.5237998	FALSE	FALSE
## magnet_belt_x	1.090141	1.6664968	FALSE	FALSE
## magnet_belt_y	1.099688	1.5187035	FALSE	FALSE
## magnet_belt_z	1.006369	2.3290184	FALSE	FALSE
## roll_arm	52.338462	13.5256345	FALSE	FALSE
## pitch_arm	87.256410	15.7323412	FALSE	FALSE
## yaw_arm	33.029126	14.6570176	FALSE	FALSE
## total_accel_arm	1.024526	0.3363572	FALSE	FALSE
## gyros_arm_x	1.015504	3.2769341	FALSE	FALSE
## gyros_arm_y	1.454369	1.9162165	FALSE	FALSE
## gyros_arm_z	1.110687	1.2638875	FALSE	FALSE
## accel_arm_x	1.017341	3.9598410	FALSE	FALSE
## accel_arm_y	1.140187	2.7367241	FALSE	FALSE
## accel_arm_z	1.128000	4.0362858	FALSE	FALSE
## magnet_arm_x	1.000000	6.8239731	FALSE	FALSE
## magnet_arm_y	1.056818	4.4439914	FALSE	FALSE
## magnet_arm_z	1.036364	6.4468454	FALSE	FALSE
## roll_forearm	11.589286	11.0895933	FALSE	FALSE
## pitch_forearm	65.983051	14.8557741	FALSE	FALSE
## yaw_forearm	15.322835	10.1467740	FALSE	FALSE
## total_accel_forearm	1.128928	0.3567424	FALSE	FALSE
## gyros_forearm_x	1.059273	1.5187035	FALSE	FALSE
## gyros_forearm_y	1.036554	3.7763735	FALSE	FALSE
## gyros_forearm_z	1.122917	1.5645704	FALSE	FALSE
## accel_forearm_x	1.126437	4.0464784	FALSE	FALSE
## accel_forearm_y	1.059406	5.1116094	FALSE	FALSE
## accel_forearm_z	1.006250	2.9558659	FALSE	FALSE
## magnet_forearm_x	1.012346	7.7667924	FALSE	FALSE
## magnet_forearm_y	1.246914	9.5403119	FALSE	FALSE
## magnet_forearm_z	1.000000	8.5771073	FALSE	FALSE

It seems like all the selected features can be informative, so we can use them to build our models.

Modeling

First of all, we need to create a training and testing subset from the `training` dataset, using respectively 70% and 30% of the starting data.

```
set.seed(420)
training$classe <- target
tridx <- createDataPartition(target, p = 0.8, list = F)
```

```
df_train <- training[tridx, ]
df_test <- training[-tridx, ]
```

Now we'll compare a couple of models. We can (hopefully) expect an error rate less than 1%. Let's start using a simple classification tree model.

```
library(rpart)
set.seed(420)
fit_tree <- train(classe ~ ., data = df_train, method = "rpart")
fit_tree
```

```
## CART
##
## 15699 samples
##    39 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 15699, 15699, 15699, 15699, 15699, 15699, ...
## Resampling results across tuning parameters:
##
##    cp          Accuracy    Kappa
## 0.02127281  0.5779991  0.46862361
## 0.03950452  0.4157213  0.21746739
## 0.11731197  0.3250498  0.06460988
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.02127281.
```

Let's see how well it works on our df_test data.

```
pred_tree <- predict(fit_tree, df_test)
confusionMatrix(pred_tree, df_test$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  A   B   C   D   E
##           A 829 197 145  86  69
##           B   2 128  12   4   5
##           C  71 246 296  69 132
##           D 193 184 231 427  75
##           E  21   4   0  57 440
##
## Overall Statistics
##
##           Accuracy : 0.5404
##           95% CI   : (0.5247, 0.5561)
##    No Information Rate : 0.2845
##    P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa   : 0.4178
##  McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
```

```
##
##               Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.7428  0.16864  0.43275  0.6641  0.6103
## Specificity      0.8229  0.99273  0.84007  0.7918  0.9744
## Pos Pred Value   0.6252  0.84768  0.36364  0.3847  0.8429
## Neg Pred Value   0.8895  0.83271  0.87520  0.9232  0.9174
## Prevalence       0.2845  0.19347  0.17436  0.1639  0.1838
## Detection Rate   0.2113  0.03263  0.07545  0.1088  0.1122
## Detection Prevalence 0.3380  0.03849  0.20749  0.2829  0.1331
## Balanced Accuracy 0.7829  0.58069  0.63641  0.7279  0.7923
```

An accuracy of 0.54 is not very promising. We can try with a random forest model instead.

```
library(randomForest)
set.seed(420)
fit_rf <- randomForest(classe ~ ., data = df_train)
fit_rf
```

```
##
## Call:
## randomForest(formula = classe ~ ., data = df_train)
##               Type of random forest: classification
##               Number of trees: 500
## No. of variables tried at each split: 6
##
## OOB estimate of error rate: 0.61%
## Confusion matrix:
##      A    B    C    D    E class.error
## A 4455     6     3     0     0 0.002016129
## B  10 3021     7     0     0 0.005595787
## C    1   19 2705    12     1 0.012052593
## D    1    0   22 2545     5 0.010882239
## E    0    0    0    8 2878 0.002772003
```

The randomForest already takes care of cross validation, and we can see an error rate of 0.61%.

Let's check its performance on the df_test subset.

```
pred_rf <- predict(fit_rf, df_test)
confusionMatrix(pred_rf, df_test$classe)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction    A    B    C    D    E
##      A 1116     0     0     0     0
##      B    0  757     5     0     0
##      C    0    2  672     6     0
##      D    0    0    7  636     0
##      E    0    0    0    1  721
##
## Overall Statistics
##
##              Accuracy : 0.9946
##              95% CI : (0.9918, 0.9967)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
```


```
##
##           Kappa : 0.9932
## Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000  0.9974  0.9825  0.9891  1.0000
## Specificity      1.0000  0.9984  0.9975  0.9979  0.9997
## Pos Pred Value   1.0000  0.9934  0.9882  0.9891  0.9986
## Neg Pred Value   1.0000  0.9994  0.9963  0.9979  1.0000
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2845  0.1930  0.1713  0.1621  0.1838
## Detection Prevalence 0.2845  0.1942  0.1733  0.1639  0.1840
## Balanced Accuracy 1.0000  0.9979  0.9900  0.9935  0.9998
```

With an accuracy of 0.9946, we can expect an out-of-sample error of 0.54%. We can safely conclude that this will be our model of choice for the rest of the project.

We might be curious about the most important features chosen by our model. Let's see the top 10.

```
varImpPlot(fit_rf, n.var = 10, main = "Random forest feature importance")
```

Random forest feature importance



Final prediction

Let's first retrain our random forest model on the whole training dataset, so we can use it for the actual prediction.

```
set.seed(420)
fit <- randomForest(classe ~ ., data = training)
```

Now we can use the trained model to predict the `testing` data.

```
pred <- predict(fit, testing)
# Actual answer not shown ;)
# pred
```